Can AI Predict ICU Outcomes Better Than Doctors' Scoring Systems?

Mohamed Gueye   Ilyes Hamzaoui

UFR des Sciences et Techniques — Université Le Havre Normandie

The Core Question

Every year, millions of patients are admitted to Intensive Care Units (ICUs). Clinicians must answer two critical questions **within hours of admission**:

Will this patient survive?

How long will they stay?

Getting these wrong means wasted resources, delayed care, and **preventable deaths**. Traditional scoring systems (APACHE, SOFA) have been the standard for decades — but they were built on linear assumptions and a handful of variables.

Machine Learning can now do better. Here's the evidence.

Why Traditional Scores Are Not Enough

Tools like **APACHE II** and **SOFA** score patients using a fixed set of rules designed in the 1980s–90s. Their limitations are well-documented:

- Assume **linear** relationships between variables and outcome
- Use only **12–17 handpicked** clinical parameters
- Calibrated on old populations — **poor generalization** today
- APACHE II AUC ≈ 0.73, SOFA AUC ≈ 0.71

In a modern ICU, EHRs capture **hundreds of variables per patient per hour**. Traditional scores leave most of this information on the table.

The Data: MIMIC-IV

Most studies — including the strongest results — use **MIMIC-IV**, the world's largest publicly available critical care database:

67,748 ICU patients | Beth Israel Deaconess Medical Center
2008–2019 | Vital signs, labs, medications, nursing notes

Features extracted from the **first 24 hours** of admission:

- Vital signs: heart rate, MAP, SpO_2, respiratory rate, temperature
- Lab values: lactate, creatinine, bilirubin, glucose, WBC
- Neurological: Glasgow Coma Scale (GCS)
- Demographics: age, sex, admission type, comorbidities

Missing data — very common in ICU records — handled with **MICE** (Multiple Imputation by Chained Equations).

Models Tested

Four major approaches have been benchmarked across studies:

- **Logistic Regression** — classic interpretable baseline
- **Random Forest** — robust ensemble, handles noise well
- **XGBoost** — gradient boosting, top performer on tabular data
- **LSTM / GRU** — recurrent networks that model *time-evolving* physiology

Unlike static models, LSTM reads the **full trajectory** of a patient's vitals over time — just like an experienced clinician would.

Results: How Much Better Is ML?

Model	Mortality AUC	vs. APACHE II
APACHE II (baseline)	0.73	—
Logistic Regression	0.75–0.78	+3–7%
Random Forest	0.82–0.86	+12–18%
XGBoost	0.87–**0.918**	+19–26%
LSTM	**0.88–0.92**	**+20–26%**

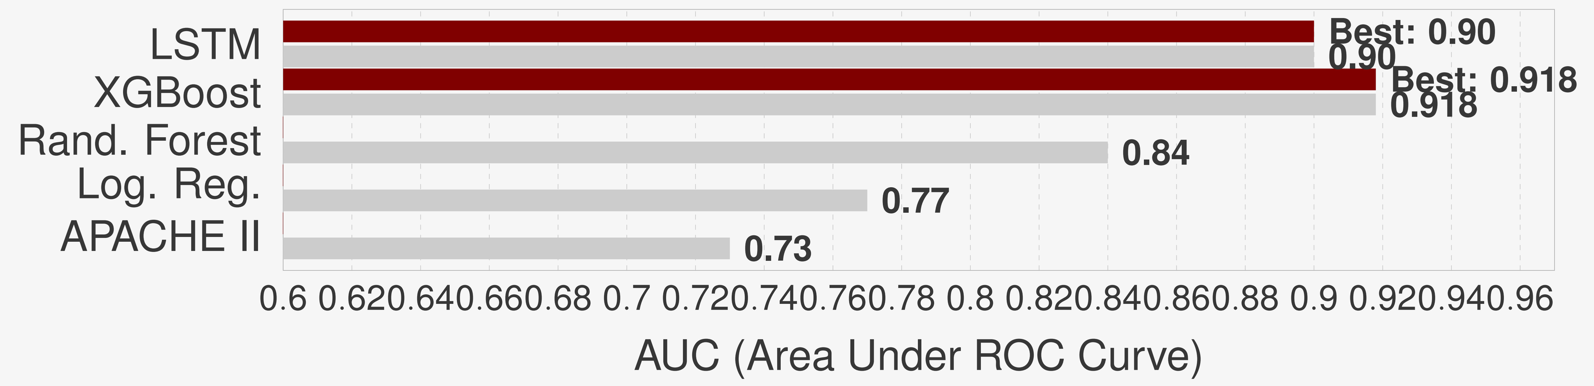

Sources: MIMIC-IV studies, 2022–2024. XGBoost AUC = 0.918 (95% CI: 0.915–0.922).

Key Takeaway

XGBoost on MIMIC-IV (67,748 patients) achieved AUC = 0.918 — a 26% improvement over APACHE II.

This means **1 in 4 high-risk patients misclassified by APACHE II** would be correctly identified by ML — enabling earlier intervention and potentially saving lives.

What Drives the Predictions? (SHAP Analysis)

Across all models, SHAP values consistently rank the same features as most predictive:

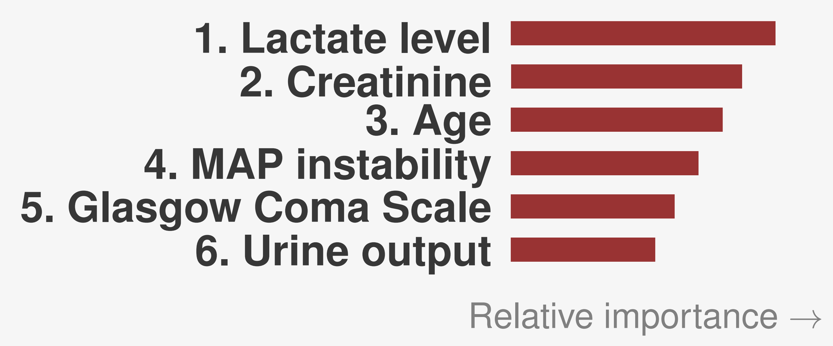

These match clinical knowledge — which is **crucial for clinician trust** and regulatory approval of AI tools in hospitals.

What About Length of Stay?

LOS prediction is **harder** than mortality — and that's expected:

- Discharge decisions involve **non-medical** factors:
  bed availability, family readiness, weekday vs. weekend staffing
- Physiological data alone cannot capture these signals
- Best models achieve MAE of 1.5–2.5 days on average ICU stays

Integrating **operational hospital data** (ward load, discharge coordinator availability) is a promising direction to close this gap.

Limitations & What Comes Next

ML models are powerful — but not yet ready for blind clinical deployment:

- Most studies are **retrospective** (historical data only)
- Models trained at one hospital **often fail to generalize** elsewhere
- **Black-box predictions** erode clinician trust
- Risk of **automation bias** — over-relying on AI scores

The roadmap to clinical deployment:

1. Prospective multi-center trials
2. Explainability tools (SHAP, LIME) integrated into EHR dashboards
3. Federated learning across hospital networks
4. AI as a *second opinion* — not a replacement for clinicians

Bottom Line

mohamed.gueye@etu.univ-lehavre.fr | ilyes.hamzaoui@etu.univ-lehavre.fr